

Before getting into what I'd change in this specific skill, I want to explain how I think about evaluating an AI workflow in general, since that's really what this exercise is testing.

The first thing I look at is where the line falls between what the model should decide and what should just be code. A calculation, a lookup, a validation rule, anything with one right answer every time doesn't belong in a paragraph of instructions for an LLM to execute mentally. It belongs in code that runs the same way on every input. The model earns its keep on the parts that actually benefit from language tone, framing, how something is phrased. A lot of the AI workflow failures I've seen come from blurring that line asking the model to do arithmetic in its head when a script would get it right every time.

Second, I look at what happens when the input isn't clean. It's easy to write a spec that only describes the happy path, but the real test of a design is what it does with a record that's missing a field, or a value that doesn't fit the expected pattern. If the answer to "what happens then" is "the model will figure it out," that's not really a design decision, it's an unaddressed gap someone downstream is going to discover the hard way.

Third, and this one matters a lot for anything a customer or donor is facing. I look for places where the model has been given authority to state something as fact to a third party, versus just drafting language for a human to review. Those are different things. An AI that writes "we're so grateful for your support" is drafting. An AI that writes "your gift will be matched" is asserting something on the organization's behalf, and that needs to actually be true.

Fourth, I want to be able to see why the system produced a given output, not just trust that it did. If a person can't quickly check the model's math or reasoning, the system isn't really trustworthy, it's just untested. Showing the work is what turns a black box into something a human can actually sign off on.

And finally, I pay attention to when a capability should not activate. Knowing the boundaries of a tool is as much a design choice as knowing what it does inside those boundaries. A trigger condition that's too broad means the tool shows up uninvited in tasks it was never meant for.

Charity-donor-outreach skill

The skill as drafted has real strengths. The logic is sound and personalized, scalable donor outreach is a genuinely good use case for AI. But it runs into most of the issues above in specific, fixable ways.

The ask-amount math was left to the model

The formula for calculating what to ask a donor for percentage of largest gift, rounding, a loyalty uplift, a volunteer bonus, a campaign multiplier was written as prose for an LLM to work through mentally. Across a donor list of any real size, that's inconsistent by construction: the same record could produce a slightly different number depending on how the model parsed the

instructions that day. I moved this into an explicit formula with a fixed order of operations, and added a sanity check so the calculated task can never exceed a donor's lifetime giving, which catches a certain class of arithmetic error before it reaches a letter.

Missing data had no real handling

The original instruction "if a field is missing, 'make reasonable assumptions and proceed'" hands the model a blank check on incomplete data. For a task that quotes real dollar figures and giving history back to a donor, that's probably the single riskiest sentence in the document. I replaced it with an explicit validation step: incomplete records get pulled out of the batch and routed to a review list with the specific field that's missing, instead of being filled in with a guess. In the working prototype, this caught all five Platinum donors, none of whom have a relationship manager on file in the source data at all so instead of inventing a name for each one, the skill now correctly says "I don't have this" and hands it to a person.

The model could tell donors something untrue

The Emergency Appeal messaging told the model to say a donor's gift would be matched, even when no match had actually been confirmed. This isn't a copy writing, it's a false statement made to induce a donation, which is a real legal and reputational problem for a real charity. The fix requires the actual confirmed match terms to be passed in as an input; if there's no match, the letter simply doesn't mention one. Of everything in this review, this was the one I'd treat as non-negotiable rather than a nice-to-have.

Nobody could check the math after the fact

The original skill hands back a finished letter with no visibility into how the ask amount or tone was decided. If a staff member wanted to sanity-check a number, they'd have to redo the entire calculation by hand. The prototype now outputs a calculation trace alongside every letter's largest gift, tier rate, adjustments, and final number so a reviewer can spot-check it in seconds instead of taking it on faith.

The skill would fire on almost anything

As written, the trigger condition activates on nearly any mention of donors, money, email, or nonprofits, which means it would try to insert donor-letter logic into something like a grant report or a volunteer schedule. I narrowed it to the actual situation the skill is built for: a donor data file plus an explicit request to generate outreach letters.

One more thing worth flagging

The donor data itself was hardcoded directly into the skill file as a static table, which cuts against the stated goal of scaling to a growing donor list every new or updated donor would mean editing the skill's source instructions rather than just updating a data file. The rewritten skill reads donor data from an external file at run time instead.

Tradeoffs I'm making

None of the fixes above are free, and I'd rather name the costs than pretend they don't exist.

Routing incomplete records to manual review means fewer letters get generated automatically. In the prototype run, five of forty-nine donors for every Platinum donor, since the source data has no relationship-manager field anywhere, didn't get an auto-generated letter. That's the right tradeoff for a workflow that quotes real dollar figures, but it's a real cost: someone has to close

that gap by hand, and in this case it means collecting an entirely new piece of data for the whole Platinum segment, not just patching one record.

Requiring a confirmed match before mentioning one probably costs something in response rate. Matching-gift language is a proven fundraising lever, and this fix means a campaign simply won't use it unless a match genuinely exists. I think that's clearly the right call, but it's worth being honest that it's not free.

And not everything I changed carries the same weight. The fabrication fix and the validation fix are correctness issues I'd defend both without much hesitation. Shifting the tone for lapsed donors from apologetic to a warmer "we'd love to have you back" framing is a different kind of change. That's a stylistic call, not a fix for something objectively broken, and it's the one item here I'd actually want fundraising staff to weigh in on rather than treating with the same certainty as everything above it.

Prototype evidence

To make sure the fixes actually hold up I built a working Python implementation of the improved skill logic and ran it against exactly the donor data given in the case study materials the same 49 donors, same fields, nothing added or removed. No synthetic donors, no extra columns. Anywhere the letter needed a fact the case study didn't supply the output shows an explicit placeholder rather than an invented value, so it's visible in the letter itself.

The original skill's Platinum-tier instructions call for "a personal relationship manager name," but no such field exists anywhere in the donor data provided. Rather than inventing a name which is what the original skill's "make reasonable assumptions" instruction would produce the rewritten skill treats this as missing required information and routes all 5 Platinum donors to manual review instead of guessing

Donor	Reason
Robert Svensson	Platinum donor has no relationship_manager on file required, not invented
Earl Fontaine	Platinum donor has no relationship_manager on file required, not invented
Walter Adeyemi	Platinum donor has no relationship_manager on file required, not invented
Ralph Osei-Bonsu	Platinum donor has no relationship_manager on file required, not invented
Victor Ambrosius	Platinum donor has no relationship_manager on file required, not invented

This is a case where the fix surfaces a real gap in the source data itself, not just a bug in the skill the tier rule references information the donor list was never structured to include. All 5 were excluded from the letter batch rather than sent with a fabricated name.

The full gift history is actually used, not just carried along

The case study table includes a year-by-year gift history for every donor. It isn't needed for the ask-amount math or the thank-you line but the skill calls for it: Annual Fund messaging should mention a donor's giving streak "only if the data shows consecutive years of giving." The prototype parses each donor's actual gift years and only adds a streak line when the years are genuinely consecutive through their most recent gift

Donor	Gift years on file	Streak line included?
Margaret Alcott	2019, 2020, 2021, 2022, 2023	Yes — 5 years running
Dorothy Callahan	2020, 2021, 2022, 2023	Yes — 4 years running
Angela Petersen	2017, 2019, 2021 (gaps)	No — correctly omitted
James Whitfield	2015, 2017, 2019, 2022 (gaps)	No — correctly omitted

Every ask amount comes with a visible calculation trace

Donor	Tier	Ask	Calculation shown
James Whitfield	Gold	\$3,100 \$3,100	$\$12,000 \times 25\% = \$3,000 \rightarrow \text{volunteer} + \$100 \rightarrow$
Dorothy Callahan	Gold	\$900	$\$3,500 \times 25\% = \$875 \rightarrow \text{rounded } \900
Margaret Alcott	Silver	\$200	$\$1,500 \times 15\% = \$225 \rightarrow \text{rounded } \200

Harold Mensah	Silver	\$300	$\$1,400 \times 15\% = \$210 \rightarrow$ rounded \$200 $\rightarrow$ volunteer +\$100 $\rightarrow$ \$300
---------------	--------	-------	--

Arthur Mwangi	Bronze	\$150	Bronze flat ask
---------------	--------	-------	-----------------